

Towards verifying data science software

Guruprerana Shabadi

Advised by **Caterina Urban**, INRIA Paris and worked with **Luca Negrini** and **Pietro Ferrara**
(Università Ca' Foscari Venezia)

Motivation

- **Data science:** Analysing and transforming raw and potentially dirty *tables of data*
- **Tedious process:** Written as **one-off** code; rarely tested
- Errors in processing data propagate into unsound analyses and inaccurate ML models

```
In [ ]: import pandas as pd

In [ ]: from fugue_notebook import setup
        setup()

In [ ]: df = pd.DataFrame({'a':[1,2,3,4], 'b':[1,2,3,4]})
        df.to_csv('df.csv', index=False)

In [ ]: %%sql
        -- This SQL cell sees the dataframe defined in the previous cell
        SELECT *
        FROM df
        WHERE a > 2
        PRINT

In [ ]: %%sql
        df2 = LOAD "/Users/kevinkho/Work/fugue/df.csv" (header=TRUE, infer_schema=TRUE)
        SELECT *
        FROM df2
        WHERE b < 2
        PRINT
        SAVE OVERWRITE "/Users/kevinkho/Work/fugue/df.csv"

In [ ]:
```



Jupyter Notebook

Goal

Statically analyse Jupyter Notebooks

(focussing on the pandas library)

Problem

Are there any **erroneous** data transforms?

```
df = pd.read('covid.csv')  
df.drop('spo2')  
  
:  
  
print(df['spo2'].mean())
```

Problem

Is there any **unused** data?
Was any **sensitive** data used?

```
# risk 'D' unused  
  
df['risk'].map(  
    {'A': 5, 'B': 4, 'C': 3})  
  
:  
  
# sensitive data used  
  
train(df['gender', 'risk'])
```

Problem

Do any transforms introduce
bias/skew?

```
df.filter('age' > 40)
```

```
:
```

```
# possible dependence  
between 'age' and 'risk'
```

```
train(df['age', 'risk'])
```

Challenges

- These problems **do not always throw errors**
- Hard to debug manually
- Notebook semantics add a level of complexity
- **Structure of the data** required to reason about data usage, bias/skew



**Infer structure and
shape of data from the
code!**

Coder places assumptions on shape of data

```
df1 = pd.read('covid1.csv')  
df1.filter('spo2' < 95)  
df1['risk'].map(  
    {'A': 5, 'B': 4, 'C': 3})
```

```
df2 = pd.read('covid2.csv')  
df2.select_rows(500, 1000)
```

```
df3 = pd.concat([df1, df2])
```

Input data shape inference

- Constraints are placed by each **dataframe transformation**
- We use **abstract interpretation** to gather all constraints placed by the code
- Column and row info, types of data in columns, bounds on values...

Dataframe graph abstract domain

Helps track all the transformations applied to dataframes

Graph construction

df1


READ

```
➔ df1 = pd.read('covid1.csv')  
df1.filter('spo2' < 95)  
df1['risk'].map(  
    {'A': 5, 'B': 4, 'C': 3})
```

```
df2 = pd.read('covid2.csv')  
df2.select_rows(500, 1000)
```

```
df3 = pd.concat([df1, df2])
```

Graph construction




```
df1 = pd.read('covid1.csv')  
➔ df1.filter('spo2' < 95)  
df1['risk'].map(  
    {'A': 5, 'B': 4, 'C': 3})
```

```
df2 = pd.read('covid2.csv')  
df2.select_rows(500, 1000)
```

```
df3 = pd.concat([df1, df2])
```

Graph construction



```
df1 = pd.read('covid1.csv')
```

```
df1.filter('spo2' < 95)
```


```
df1['risk'].map(  
    {'A': 5, 'B': 4, 'C': 3})
```

```
df2 = pd.read('covid2.csv')
```

```
➔ df2.select_rows(500, 1000)
```

```
df3 = pd.concat([df1, df2])
```

Graph construction




```
df1 = pd.read('covid1.csv')  
df1.filter('spo2' < 95)  
df1['risk'].map(  
    {'A': 5, 'B': 4, 'C': 3})
```

```
df2 = pd.read('covid2.csv')  
df2.select_rows(500, 1000)
```

➔ `df3 = pd.concat([df1, df2])`

Graph construction




```
df1 = pd.read('covid1.csv')  
df1.filter('spo2' < 95)  
df1['risk'].map(  
    {'A': 5, 'B': 4, 'C': 3})
```

```
df2 = pd.read('covid2.csv')  
df2.select_rows(500, 1000)
```

➔ `df3 = pd.concat([df1, df2])`

Conditionals

(least upper bound)



```
df = pd.read('covid1.csv')
```

```
if 'risk' in df:
```

```
    df['risk'].map(map1)
```

```
else:
```

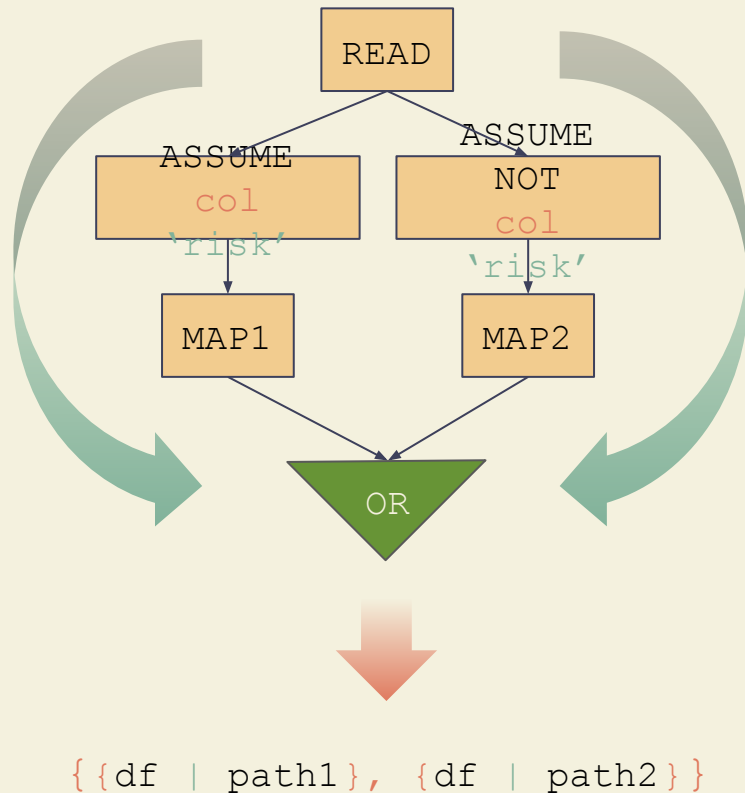
```
    df['risk'].map(map2)
```


Dataframe graphs

Graph node = **Constraint**

Graph = **Set of set of
Dataframes**


One set corresponding to
each *OR-path* of the graph






Without OR-nodes, we only have one set of dataframes

Partial order on dataframe graphs



`{{df | path1}}`



`{{df | path1}, {df | path2}}`



**The set of dataframe
graphs forms a lattice!**


Dealing with loops using widening

- Loops would add unbounded nodes to graphs
- We use **naive widening** to avoid this, i.e., jump to \top
- Alternative: introduce **cycles!**

```
# luckily this is uncommon  
while bool_cond:  
    Op1(df)
```

Using dataframe graphs

1. **Restrict** to condensed set of dataframe transformations
2. **Understand** the constraints placed by each transform
3. **Collect** all constraints on input data files from all the graphs



Implementation

- The dataframe graph abstract domain was implemented as a part of PyLiSA
- Built on LiSA - a modular multi-language static analyser
- We can produce graphs from notebooks for a subset of pandas
- Able to specify order of execution of cells within notebooks

Next step

**Use the generated graphs to
reason about input data
usage, data provenance,
bias inducing transforms**

Thank you!